

dev1num2

February 5, 2019

```
In [8]: import torch
import torchvision
import torchvision.transforms as transforms #
import torchvision.datasets.mnist as mnist # to import data

# we use torch.cuda.Event(enable_timing=True) to measure time
# from timeit import default_timer as timer
# import time

import torch.optim as optim
import torch.nn as nn
import torch.nn.functional as F

import collections          # for ordered_dictionnary
import torch.nn.init as init # to initialize model

import copy                  # for copy.deepcopy( ... )

import matplotlib.pyplot as plt
import numpy as np

import collections
import torch.nn.init as init
```

1 Assignment 1, part 2

1.1 import data

```
In [3]: transform1 = transforms.Compose( [transforms.ToTensor(), transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))] )
transform2 = transforms.Compose( [transforms.ToTensor(), transforms.Normalize( (0.1307, 0.1307, 0.1307), (0.3081, 0.3081, 0.3081))] )
dl_path = "./data_mnist"

trainset = mnist.MNIST( root = dl_path, train=True , download=True, transform = transform1 )
testset = mnist.MNIST( root = dl_path, train=False, download=True, transform = transform2 )
```

display some samples

```
In [27]: nb_sample = 8
```

```
trainloader = torch.utils.data.DataLoader(trainset, batch_size = nb_sample, shuffle=True)
testloader = torch.utils.data.DataLoader(testset, batch_size = nb_sample, shuffle=False)
```

```
# functions to show an image
```

```
def imshow(img):
    img = ( img + 0.43 ) / 3.3    # to be in the interval [0,1]
    npimg = img.numpy()
    npimg = (255*npimg).astype(np.uint8) # to be a int in (0,...,255)
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()
```

```
# get some random training images
```

```
dataiter = iter(trainloader)
images, labels = dataiter.next()
```


```
# show images
```

```
imshow(torchvision.utils.make_grid(images))
```

```
# print labels
```

```
print(' '.join('%5s' % labels[j].item() for j in range(nb_sample)))
```

```
print( "img size = " , images.size() , "label size = " , labels.size() )
```



```
7      2      3      3      4      7      4      8
img size = torch.Size([8, 1, 28, 28]) label size = torch.Size([8])
```

```
print size of train-test dataset
```

```
In [5]: training_dataset_size = trainset.__len__()
        testing_dataset_size = testset.__len__()
```

```
print( training_dataset_size )
```

```
print( testing_dataset_size )
```

```
60000
```

```
10000
```

set the device to cuda if possible

```
In [6]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
        print(device)
```

cuda:0

1.2 Define some modules

1.2.1 the MLP

```
In [9]: class MLP(nn.Module):
        def __init__(self, h1, h2):
            super(MLP, self).__init__()
            self.fc1 = nn.Linear(784, h1)
            self.fc2 = nn.Linear(h1, h2)
            self.fc3 = nn.Linear(h2, 10)

        def forward(self, x):
            x = x.view(-1, 28 * 28)
            x = F.relu(self.fc1(x))
            x = F.relu(self.fc2(x))
            x = F.softmax(self.fc3(x), dim=-1)
            return x
```

```
In [10]: # fan_in, fan_out = _calculate_fan_in_and_fan_out(tensor)
        # print( nn.init._calculate_fan_in_and_fan_out(net.state_dict()["fc1.weight"]) )
        # print( net.state_dict()["fc1.weight"].size() )
```

1.2.2 The CNN

architecture taken from : https://github.com/MaximumEntropy/welcome_tutorials/tree/pytorch/pytorch

```
In [11]: class Classifier(nn.Module):
        """Convnet Classifier"""
        def __init__(self, kernel_sz = 3 ):

            if kernel_sz % 2 == 0 :
                raise ValueError("kernel size must be odd")
            pad = kernel_sz // 2

            super(Classifier, self).__init__()
            self.conv = nn.Sequential(
                # Layer 1
                nn.Conv2d(in_channels=1, out_channels=16, kernel_size=( kernel_sz , kernel_sz ),
                           padding=pad),
                # nn.Dropout(p=0.5),
                nn.ReLU(),
                nn.MaxPool2d(kernel_size=(2, 2), stride=2),
```

```

        # Layer 2
        nn.Conv2d(in_channels=16, out_channels=32, kernel_size=( kernel_sz , kernel_sz ),
        # nn.Dropout(p=0.5),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=(2, 2), stride=2),

        # Layer 3
        nn.Conv2d(in_channels=32, out_channels=64, kernel_size=( kernel_sz , kernel_sz ),
        # nn.Dropout(p=0.5),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=(2, 2), stride=2),

        # Layer 4
        nn.Conv2d(in_channels=64, out_channels=128, kernel_size=(kernel_sz , kernel_sz ),
        # nn.Dropout(p=0.5),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=(2, 2), stride=2)
    )
    # Logistic Regression
    self.clf = nn.Linear(128, 10)

def forward(self, x):
    return self.clf(self.conv(x).squeeze())

```

Print the number of parameters in each models and display the computation

```

In [12]: def number_of_params( net ) :
    nb_param = 0
    param_lst = " "
    for i, (key, value) in enumerate( net.state_dict().items() ) :

        if i == 0 :
            param_lst = param_lst + "\n ({{:<20}}      ".format( key + " )" )
        else :
            param_lst = param_lst + "\n ({{:<20}}  + ".format( key + " )" )

    nb_param_tmp = 1

    for j , x in enumerate(value.size()) :
        if j == 0 :
            param_lst = param_lst + "{{xx}}".format( xx = x )
        else :
            param_lst = param_lst + "*{{xx}}".format( xx = x )

    nb_param_tmp = nb_param_tmp * x

```

```

nb_param = nb_param + nb_param_tmp

print( "number of params = " , nb_param , " = " , param_lst )

mytestnet1 = MLP(h1=620,h2=620)
mytestnet2 = Classifier( kernel_sz=9 )
number_of_params( mytestnet1 )
number_of_params( mytestnet2 )
del mytestnet1
del mytestnet2

number of params = 877930 =
(fc1.weight)          620*784
(fc1.bias)             + 620
(fc2.weight)          + 620*620
(fc2.bias)             + 620
(fc3.weight)          + 10*620
(fc3.bias)             + 10
number of params = 873738 =
(conv.0.weight)        16*1*9*9
(conv.0.bias)          + 16
(conv.3.weight)        + 32*16*9*9
(conv.3.bias)          + 32
(conv.6.weight)        + 64*32*9*9
(conv.6.bias)          + 64
(conv.9.weight)        + 128*64*9*9
(conv.9.bias)          + 128
(clf.weight)           + 10*128
(clf.bias)             + 10

```

1.2.3 Define some initialization methods

```

In [13]: def gloriot_init ( layer ) :
         """
         Weiths are generated from  $U[-d,d]$  where  $d = \sqrt{6/(fan\_in + fan\_out)}$ , biases are
         """
         if type(layer) == nn.Linear or type(layer) == nn.Conv2d :
             init.xavier_uniform_( layer.weight , gain=1 )
             layer.bias.data.fill_(0.0)

         def zero_init ( layer ) :
             """Everything is set to zero"""
             if type(layer) == nn.Linear or type(layer) == nn.Conv2d :
                 layer.weight.data.fill_(0.0)
                 layer.bias.data.fill_(0.0)

         def norm_init ( layer ) :

```

```

        """Weights are generated from std normal, biases are set to zero"""
        if type(layer) == nn.Linear or type(layer) == nn.Conv2d :
            init.normal_(layer.weight, mean=0, std=1)
            layer.bias.data.fill_(0.0)

In [14]: # net = MLP(500,500)
        # print( net.state_dict()["fc1.weight"] )
        # print( net.state_dict()["fc1.bias"] )
        # net.apply( one_step_Glorot_init )
        # print( net.state_dict()["fc1.weight"] )
        # print( net.state_dict()["fc1.bias"] )

In [100]: if False :
        del cudanet

In [15]: cudanet = Classifier( kernel_sz = 9 )
        # cudanet = MLP(h1=620,h2=620)

        cudanet.apply( glorot_init )
        # cudanet.apply( zero_init )
        # cudanet.apply( norm_init )

        # cudanet.load_state_dict(torch.load(path), strict=False)
        cudanet.to(device)

Out[15]: Classifier(
  (conv): Sequential(
    (0): Conv2d(1, 16, kernel_size=(9, 9), stride=(1, 1), padding=(4, 4))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(16, 32, kernel_size=(9, 9), stride=(1, 1), padding=(4, 4))
    (4): ReLU()
    (5): MaxPool2d(kernel_size=(2, 2), stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(32, 64, kernel_size=(9, 9), stride=(1, 1), padding=(4, 4))
    (7): ReLU()
    (8): MaxPool2d(kernel_size=(2, 2), stride=2, padding=0, dilation=1, ceil_mode=False)
    (9): Conv2d(64, 128, kernel_size=(9, 9), stride=(1, 1), padding=(4, 4))
    (10): ReLU()
    (11): MaxPool2d(kernel_size=(2, 2), stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (clf): Linear(in_features=128, out_features=10, bias=True)
)

In [16]: criterion = nn.CrossEntropyLoss()
        optimizer = optim.SGD(cudanet.parameters(), lr=0.055, momentum=0.0, weight_decay=0)
        nb_epoch = 10

In [17]: trainloader = torch.utils.data.DataLoader(trainset, batch_size=16*64,shuffle=True, num
        state_dict_list = list()

```

```

torch.cuda.synchronize()

start = torch.cuda.Event(enable_timing=True)
end   = torch.cuda.Event(enable_timing=True)
start.record()

for epoch in range( nb_epoch ):  # loop over the dataset multiple times

    # if epoch == 5 :
    #     optimizer = optim.SGD(cudanet.parameters(), lr=0.05, momentum=0.0, weight_decay=0.0001)

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = cudanet(inputs)
        # if i == 0 : print( outputs )
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
    else : # print every epoch
        print('epoch = %d, loss = %.8f' % (epoch + 1, running_loss / (i*8*64))) # nb of iterations
        running_loss = 0.0
        torch.cuda.synchronize()
        tmp_state_dict = {}
        for k, v in cudanet.state_dict().items():
            tmp_state_dict[k] = v.cpu()
        state_dict_list.append( tmp_state_dict )
        torch.cuda.synchronize()

    else :
        print('Finished Training')

    end.record()
    torch.cuda.synchronize()
    print( "time required = " , start.elapsed_time(end)*0.001 , " s ")

```

epoch = 1, loss = 0.00347052

```

epoch = 2, loss = 0.00084709
epoch = 3, loss = 0.00039708
epoch = 4, loss = 0.00031282
epoch = 5, loss = 0.00021171
epoch = 6, loss = 0.00018068
epoch = 7, loss = 0.00015818
epoch = 8, loss = 0.00014084
epoch = 9, loss = 0.00012693
epoch = 10, loss = 0.00011423
Finished Training
time required = 139.8591875 s

```

test accuracy of net in its current state on the validation set

```

In [75]: testloader = torch.utils.data.DataLoader(testset, batch_size=8*64,shuffle=True, num_w
correct = torch.tensor([0])
total = torch.tensor([0])

correct, total = correct.to(device) , total.to(device)
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = cudanet(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        # print(correct)
        # print(predicted.size())
        # print(labels.size())
        # break
        correct += (predicted == labels).sum()

print('Accuracy of the network on the' , testing_dataset_size , 'test images: %.2f %%'
      % ( (100 * correct.double()) / total.double() )
)

```

Accuracy of the network on the 10000 test images: 91.18 %

Display the current state of the net

```

In [77]: if False :
        for key, value in cudanet.state_dict().items() :
            pass
            print( key , " = \n" , value , end = "\n")

```

training accuracy


```

In [96]: trainloader = torch.utils.data.DataLoader(trainset, batch_size=16*64,shuffle=True, num

correct = torch.tensor([0])
total = torch.tensor([0])

correct, total = correct.to(device) , total.to(device)
with torch.no_grad():
    for data in trainloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = cudanet(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        # print(correct)
        # print(predicted)
        # print(labels)
        correct += (predicted == labels).sum()

print('Accuracy of the network on the' , training_dataset_size , 'train images: %.2f
      % ( (100.0 * correct.double() ) / total.double() )
      )

```

Accuracy of the network on the 60000 train images: 98.32 %

compare accuracy across epoch on training-test dataset

```

In [24]: trainloader = torch.utils.data.DataLoader(trainset, batch_size = 8*64, shuffle=True ,
testloader  = torch.utils.data.DataLoader(testset , batch_size = 8*64, shuffle=True ,

accuracy = torch.ones(nb_epoch,2, dtype=torch.float) * 100
for epoch , tmp_state_dict in enumerate(state_dict_list,0) :
    if epoch % 1 != 0 :
        continue

    # net_model = cudanet.__class__
    # cudaTOCPUnet = net_model(h1,h2)
    # cudaTOCPUnet = Classifier( kernel_sz = 9 )
    cudaTOCPUnet = copy.deepcopy( cudanet )
    cudaTOCPUnet.load_state_dict( tmp_state_dict )
    cuda_test_net = copy.deepcopy(cudaTOCPUnet).to(device)

    correct = torch.tensor([0,0])
    total  = torch.tensor([0,0])

    correct, total = correct.to(device) , total.to(device)

    loader_list = [ testloader , trainloader ]

```

```

with torch.no_grad():
    for i, loader in enumerate(loader_list,0) :
        for data in loader:
            images, labels = data
            images, labels = images.to(device), labels.to(device)
            outputs = cuda_test_net(images)
            _, predicted = torch.max(outputs.data, 1)
            total[i] += labels.size(0)
            correct[i] += (predicted == labels).sum()

accuracy[epoch,:] = accuracy[epoch,:] * correct.type(torch.FloatTensor) / total.t

print('epoch %3d : Accuracy of the network on the test images: %.2f %% , training
      % ( epoch+1, accuracy[epoch,0] , accuracy[epoch,1] )
    )

```

```

epoch  1 : Accuracy of the network on the test images: 92.44 % , training images 91.84 %
epoch  2 : Accuracy of the network on the test images: 95.77 % , training images 95.63 %
epoch  3 : Accuracy of the network on the test images: 96.42 % , training images 96.55 %
epoch  4 : Accuracy of the network on the test images: 96.96 % , training images 97.02 %
epoch  5 : Accuracy of the network on the test images: 97.54 % , training images 97.55 %
epoch  6 : Accuracy of the network on the test images: 97.28 % , training images 97.31 %
epoch  7 : Accuracy of the network on the test images: 97.76 % , training images 97.94 %
epoch  8 : Accuracy of the network on the test images: 98.04 % , training images 98.19 %

```

Plot the train and validation errors for each of the model

```

In [25]: from matplotlib.pyplot import figure
        # import warnings
        # warnings.filterwarnings('ignore')

        # data to plot
        n_groups = nb_epoch
        accuracy_toplot = accuracy.numpy()

        tests_accuracy = accuracy_toplot[:,0]
        train_accuracy = accuracy_toplot[:,1]

        #
        plt.rcParams.update({'font.size': 16})
        plt.rcParams["figure.figsize"] = (12 ,8)
        # create plot
        fig, ax = plt.subplots()
        index = np.arange(n_groups)
        bar_width = 0.3
        opacity = 0.8

```

```

rects1 = plt.bar(index, tests_accuracy, bar_width,
                  alpha=opacity,
                  color='g',
                  label='Test set')

rects2 = plt.bar(index + bar_width, train_accuracy, bar_width,
                  alpha=opacity,
                  color='r',
                  label='Training set')

eps = 3
top = min(int( np.ceil(accuracy_toplot.max() + eps)) , 100)
bot = max(int(np.floor(accuracy_toplot.min() - eps)) , 0 )

plt.ylim(bot, top)      # set the ylim to bottom, top
plt.axhline(y=97,color="black")

plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Comparison between training set and test set accuracy \nduring the training')
plt.xticks(index + bar_width, range(1,11,1) )
plt.yticks( range(bot,top+1,1) )
plt.legend()

# plt.tight_layout()
plt.show()

```



1.3 Some additional stuff

3-best accuracy : having the right label in the net's top 3 answers count as a good answer

```
In [58]: correct = torch.tensor([0])
         total = torch.tensor([0])

correct, total = correct.to(device) , total.to(device)
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = cudanet(images)
        predicted = torch.topk(outputs.data, 3)[1]
        total += labels.size(0)
        correct += (predicted[:,0] == labels).sum()
        correct += (predicted[:,1] == labels).sum()
        correct += (predicted[:,2] == labels).sum()

print('3-best Accuracy of the network on the ' , testing_dataset_size , ' test images
      % ( 100 * correct.double() / total.double())
      )
```


3-best Accuracy of the network on the 10000 test images: 99.920 %

print some test sample that the net misclassifies

```
In [90]: nb_of_error = torch.tensor([6])
         j = torch.tensor([0])
         errorimages = torch.empty(6,1,28,28)
         errorlabels = torch.empty(6)
         erroroutputs = torch.empty(6, 10)
         for images, labels in testloader :
             images, labels = images.to(device), labels.to(device)
             outputs = cudanet(images)
             _, predicted = torch.max(outputs.data, 1)
             if not (predicted == labels).sum() == labels.size()[0] :
                 comparison = (predicted == labels)
                 for i,b in enumerate( comparison ) :
                     if b == 0 :
                         errorimages[j,:,:,:] = copy.deepcopy(images[i,:,:,:]).cpu()
                         errorlabels[j] = labels[i].clone().detach().requires_grad_(False).cpu()
                         erroroutputs[j,:] = outputs[i,:].clone().detach().requires_grad_(False).cpu()
                         j = j + 1
                         if j.item() >= nb_of_error.item() :
                             break
                     else :
                         continue
                 break

         if j.item() == 0 :
             print( "no error found" )
         else :
             imshow( torchvision.utils.make_grid(errorimages) )

             print( "this should be : " , ",".join( "%2d" % nb.item() for nb in errorlabels ) )
             for i in range(3) :
                 print( "net choice #" , i+1 , ": " , ",".join( "%2d" % nb.item() for nb in li
```



```

this should be :    8, 4, 6, 9, 5, 6
net choice # 1 :    3, 2, 0, 7, 3, 0
net choice # 2 :    8, 4, 6, 9, 5, 6
net choice # 3 :    2, 6, 5, 8, 8, 5

```

1.3.1 Save and load models

Save the state_dict of the model for each epoch on a local directory

```

In [21]: local_path = "./save"
         for epoch , tmp_state_dict in enumerate(state_dict_list,1) :
             saving_path = local_path + "/dev1num2model_for_epoch{Epoch}.pth".format( Epoch = epoch )
             torch.save( tmp_state_dict , saving_path)

./save/dev1num2model_for_epoch1.pth
./save/dev1num2model_for_epoch2.pth
./save/dev1num2model_for_epoch3.pth
./save/dev1num2model_for_epoch4.pth
./save/dev1num2model_for_epoch5.pth
./save/dev1num2model_for_epoch6.pth
./save/dev1num2model_for_epoch7.pth
./save/dev1num2model_for_epoch8.pth
./save/dev1num2model_for_epoch9.pth
./save/dev1num2model_for_epoch10.pth

```

load from file and set the load the state_dict of the last epoch on a object the files have to be located in "./save" and named "dev1num2model_for_epoch{j}.pth" for j from ... to ...

```

In [23]: local_path = "./save"
         state_dict_list = list()

         from_idx = 1
         to_idx = 10

         for epoch in range( from_idx + 1 , to_idx , 1 ):
             path = local_path + "/dev1num2model_for_epoch{Epoch}.pth".format( Epoch = epoch )
             tmp_dict = torch.load(path)
             state_dict_list.append(tmp_dict)
         else :
             cudanet = Classifier( kernel_sz = 9 )
             cudanet.load_state_dict(tmp_dict)
             cudanet.eval()

```